

ORM & SQLAlchemy

Topics Covered

- ORM fundamentals & use cases
- SQLAlchemy architecture (Engine, Session, DBAPI)
- SQLAlchemy Core vs ORM
- MetaData, schema management & reflection
- Inspector vs MetaData

ORM Concepts & Architecture

ORM (Object Relational Mapping)

What is ORM?

- Bridging the gap between **Python objects** and **relational databases**.
- ORM maps **tables** → **classes**
- Rows → objects
- Columns → attributes

Why ORM is Used

- Faster development (less SQL writing)
- Cleaner, maintainable code
- DB-agnostic (Postgres, MySQL, SQLite)
- Built-in security (SQL injection protection)

SQL vs ORM

SQL	ORM
Full control	Faster dev
Verbose	Cleaner code
DB-specific	DB-independent
Harder to maintain	Easier to refactor

SQLAlchemy

- SQLAlchemy is a **Python SQL toolkit + ORM** that lets you work with databases using **Python code instead of raw SQL**.
- **Key Points**
 - Acts as a **bridge** between Python objects and relational databases
 - Supports **ORM + Core (SQL Expression Language)**
 - Works with PostgreSQL, MySQL, SQLite, Oracle, MSSQL
 - Gives you **ORM convenience** without hiding SQL completely

SQLAlchemy Philosophies

- **Explicit over implicit**

You see what SQL is being generated. No magic.

- **SQL is not the enemy**

SQLAlchemy respects SQL — it doesn't try to hide it.

- **Separation of concerns**

Engine (DB) ≠ Session (transactions) ≠ Models (business logic)


- **One ORM, many databases**

Same code, different DBs. Minimal pain.

- **Control is optional, not forbidden**

Start simple with ORM → drop to raw SQL when needed.

Architecture



SQLAlchemy – Core

SQLAlchemy Core is the **low-level foundation** that deals directly with databases and SQL.

- **Engine**
Entry point to the database. Manages DB connectivity.
- **Dialect**
Translates generic SQLAlchemy commands into DB-specific SQL
(Postgres, MySQL, SQLite, etc.)
- **Connection Pool**
Maintains reusable DB connections for performance.
- **SQL Expression Language**
Write SQL using Python expressions (SELECT, INSERT, JOIN).
- **Schema / Types**
Python objects represent tables, columns, and datatypes.

SQLAlchemy – ORM

SQLAlchemy ORM sits on top of the Core and provides an **object-centric view**.


- Maps **Python classes** → **database tables**
- Maps **objects** → **rows**
- Handles persistence using the **Unit of Work pattern**
- Automatically generates SQL from mappings
- Uses **Core internally** to talk to the database

Perspective shift:


- Core → schema centric
- ORM → object centric

SqlAlchemy is like Onion

Can be learned from the inside out, or outside in



Level 1: Engines, connections, transactions



Python DBAPI (PEP 249)

- Python DBAPI is a **standard specification** for how Python talks to relational databases.
- **What it Defines**
 - How to **connect** to a database
 - How to **execute SQL**
 - How to **fetch results**
- How **transactions** work (commit / rollback)
- **Common DBAPI Drivers**
 - psycopg2 → PostgreSQL
 - mysqlclient → MySQL
 - sqlite3 → SQLite
 - cx_Oracle → Oracle

DBAPI Example


```
import psycopg2
connection = psycopg2.connect("scott", "tiger", "test")

cursor = connection.cursor()
cursor.execute(
    "select emp_id, emp_name from employee "
    "where emp_id=%(emp_id)s",
    {'emp_id':5})
emp_name = cursor.fetchone()[1]
cursor.close()

cursor = connection.cursor()
cursor.execute(
    "insert into employee_of_month "
    "(emp_name) values (%(emp_name)s)",
    {"emp_name":emp_name})
cursor.close()

connection.commit()
```

SQLAlchemy Engine & DBAPI



Engine is the **core interface** between SQLAlchemy and the database.


What the Engine Does

- Manages **database connections**
- Uses **connection pooling**
- Talks to the database via **DBAPI drivers**
- Executes SQL (generated by Core or ORM)

Role of DBAPI

- DBAPI is the **actual driver** (`psycopg2`, `mysqlclient`, `sqlite3`)
- Engine **does not speak DB protocol**
- Engine delegates low-level work to DBAPI

Level 2: MetaData, Tables



MetaData (More Than Just Tables)

MetaData is the **central registry** of your database schema in SQLAlchemy.

- **What MetaData Actually Stores**

- All **Table** objects
- Column definitions & types
- **Primary keys, foreign keys**
- Constraints & indexes
- Naming conventions

What is MetaData used for?

MetaData is **SQLAlchemy's notebook** where it writes down:

- What tables exist
- What columns they have
- How tables are connected

MetaData

- **MetaData is NOT a Database**
 - It does **nothing by itself**
 - No tables are created until you call: `metadata.create_all(engine)`
- **MetaData Knows Table Order**
 - SQLAlchemy figures out **dependency order**
 - Foreign keys are created safely
 - Drop order is also handled automatically
- **Reflection is MetaData's Superpower**
 - Loads schema from an **existing DB**
 - Zero Table definitions required
 - `Table("user", metadata, autoload_with=engine)`

Multiple MetaData Objects (Why & When)

- Why create 2 MetaData objects?
 - One for **tables you define**
 - One for **tables that already exist**
 - `metadata = MetaData() # tables we define`
 - `metadata2 = MetaData() # tables we reflect`
- Reflection (reading existing DB schema)
 - No CREATE TABLE needed
 - SQLAlchemy reads structure from DB
 - Useful for **legacy databases**
- Drop ALL tables (dangerous but powerful)
 - `metadata.drop_all(engine)`

```
user_reflected = Table(  
    "user",  
    metadata2,  
    autoload_with=engine  
)
```


Inspector (Database X-Ray Tool)

- **Inspector** lets you **ask questions about the database**.
 - from sqlalchemy import inspect
 - inspector = inspect(engine)
 - → List all tables -- inspector.get_table_names()
 - → See column names, types, nullability - inspector.get_foreign_keys("address")
- **When Inspector is used**
 - Debugging schema issues
 - Reverse-engineering databases
 - Writing migration tools

Why Inspector alone is NOT enough

- **Inspector can only:**
 - List tables
 - Show columns
 - Show foreign keys
 - **You cannot:**
 - Build SQL expressions
 - Join tables
 - Insert / update rows
 - Use ORM
 - Run migrations
 - **Inspector is read-only.**
- With reflected MetaData
 - `user = Table("user", metadata, autoload_with=engine)`
 - You can now:
 - `select(user).where(user.c.id == 1)`
 - Write queries
 - Use Core / ORM
 - Participate in transactions
 - Generate migrations

Level 3: SQL Expressions




SQLAlchemy Core – SQL Expression Language

- Pythonic way to **build SQL** using objects
- Produces **database-agnostic SQL**
- Foundation for both **Core and ORM**
- **Key Ideas**
 - Table.c.column → column reference
 - == , > , in_() → SQL expressions
 - & , | , and_() , or_() → AND / OR
 - Expressions compile differently per DB dialect
- **Transactions (2.x rule)**
 - All INSERT / UPDATE / DELETE must be inside: **with engine.begin():**

```
user.c.username == "ed"  
user.c.id > 5  
user.c.username.in_(["ed", "jack"])
```

Level 4: ORM



ORM (Object Relational Mapping)


- **ORM = Bridge between Objects and Tables**
 - Maps **classes** → **tables**
 - Maps **objects** → **rows**
 - Maps **attributes** → **columns**
- You work in **code**, ORM talks **SQL**.

Why ORM Exists

- Writing raw SQL everywhere is:
 - repetitive
 - error-prone
 - hard to maintain
- Applications think in **objects**, DB thinks in **tables**
- ORM solves the **object-table mismatch**.

What ORM Actually Does

The most basic task is to translate between a domain object and a table row.



- Converts Python objects to SQL
- Executes SQL via DB drivers
- Converts result rows back to objects
- Manages transactions & sessions
- **Important:**
ORM does **not** replace SQL — it generates it.

Core Components of an ORM

- **Model** → represents a table
- **Session / Unit of Work** → manages transactions
- **Mapper** → maps class ↔ table
- **Query Builder** → builds SQL
- These pieces work together — not magic.